CC Q&A

🔰 = Important

Unit 1 -

▼ 🔰 **Q1. Define Cloud Computing. State The Characteristics and Benefits.**

- Cloud computing refers to the delivery of various computing services such as servers, storage, databases, networking, software, and more over the internet ("the cloud").
- Instead of owning their own computing infrastructure or data centers, companies can rent access to anything from applications to storage from a cloud service provider.
- Cloud computing is a model for enabling ubiquitous (being everywhere at the same time), convenient, on demand network access to a shared pool of configurable computing resources that can be rapidly provisioned and released with minimal management effort or service provider interaction.

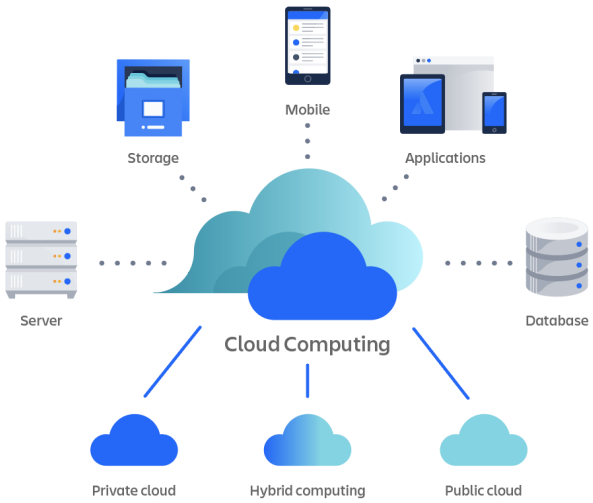

Characteristics of Cloud Computing -

1. **Ubiquitous Computing:** Cloud services are available over the network and accessible everywhere at the same time for all user.
2. **On-Demand Self-Service:** Users can provision computing resources as needed, automatically without human intervention from the service provider.
3. **Resource Pooling:** Providers serve multiple consumers with pooled computing resources, allowing for economies of scale and efficiency improvements.

4. **Rapid Elasticity:** Resources can be scaled up or down quickly and easily, allowing for flexibility in meeting varying workloads and demand spikes.

Benefits of Cloud Computing -

1. **Cost Saving:** Pay only for what is used, eliminating upfront infrastructure and maintenance costs.
2. **Accessibility:** Access services from anywhere, supporting remote work.
3. **Scalability:** Easily adjust resources to meet demand without over-provisioning.
4. **Disaster Management:** Cloud ensures data integrity and business continuity with backup and recovery solutions.
5. **Automatic Updates & Maintenance:** Providers manage updates, freeing IT from routine maintenance.

▼ 🔰 **Q2. Explain Types of Cloud Computing/Reference Model.**

- The National Institute of Standards and Technology (NIST) defines a cloud computing reference model that outlines the essential characteristics and service models of cloud computing.
- This reference model provides a comprehensive framework for understanding cloud computing services and how they are categorized and deployed.
- It ensures that both providers and users have a common understanding of the structure and capabilities of cloud services.

Types of Cloud Computing/Reference Model -

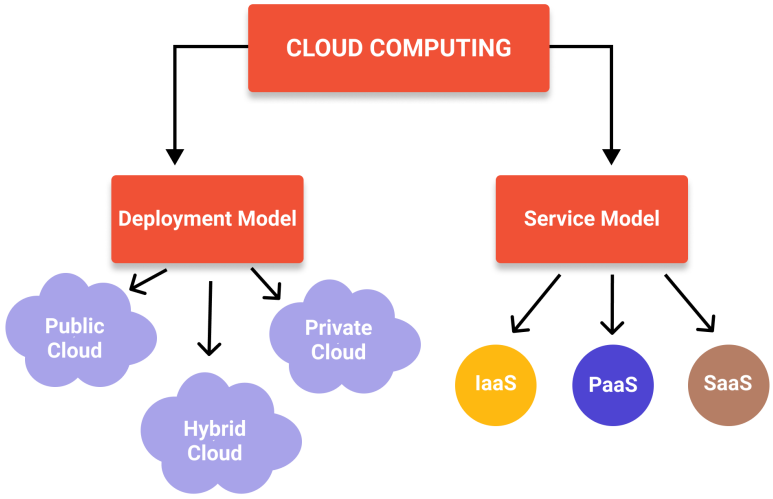

Deployment Models -

1. **Public Cloud -**
    - Public clouds deliver resources, such as compute, storage, network, develop-and-deploy environments, and applications over the internet.
    - They are owned and run by third-party cloud service provider.
    - **Example:** AWS, Microsoft Azure, Google Cloud.
2. **Private Cloud -**
    - Private clouds are built, run, and used by a single organization, typically located on-premises.
    - They provide greater control, customization, and data security but come with similar costs and resource limitations associated with traditional IT environments.
    - **Example:** VMware Cloud, OpenStack.
3. **Hybrid Cloud -**
    - Environments that mix at least one private computing environment (traditional IT infrastructure or private cloud, including edge) with one or more public clouds are called hybrid clouds.
    - They allow you to leverage the resources and services from different computing environments and choose which is the most optimal for the workloads.
    - **Example:** Microsoft Azure Hybrid Cloud, AWS Outposts.

Service Models -

1. **Infrastructure as a Service (IaaS) -**
    - IaaS delivers on-demand infrastructure resources, such as compute, storage, networking, and virtualization.
    - With IaaS, the service provider owns and operates the infrastructure, but customers will need to purchase and manage software, such as operating systems, middleware, data, and applications.
    - **Examples:** Amazon Web Services (AWS), Microsoft Azure, Google Cloud.
2. **Platform as a Service (PaaS) -**
    - PaaS delivers and manages hardware and software resources for developing, testing, delivering, and managing cloud applications. Providers typically offer middleware, development tools, and cloud databases within their PaaS offerings.
    - **Examples:** Google App Engine, Microsoft Azure App Services, Heroku.
3. **Software as a Service (SaaS) -**
    - SaaS provides a full application stack as a service that customers can access and use. SaaS solutions often come as ready-to-use applications, which are managed and maintained by the cloud service provider.

- **Examples:** Google Workspace (formerly G Suite), MS 365, Dropbox.

▼ 🔰 Q3. What Is the Difference Between Parallel and Distributed Computing?

- In **Parallel Computing** multiple processors performs multiple tasks assigned to them simultaneously. Memory in parallel systems can either be shared or distributed. Parallel computing provides concurrency and saves time and money.
- In **Distributed Computing** we have multiple autonomous computers which seems to the user as single system. In distributed systems there is no shared memory and computers communicate with each other through message passing. In distributed computing a single task is divided among different computers.

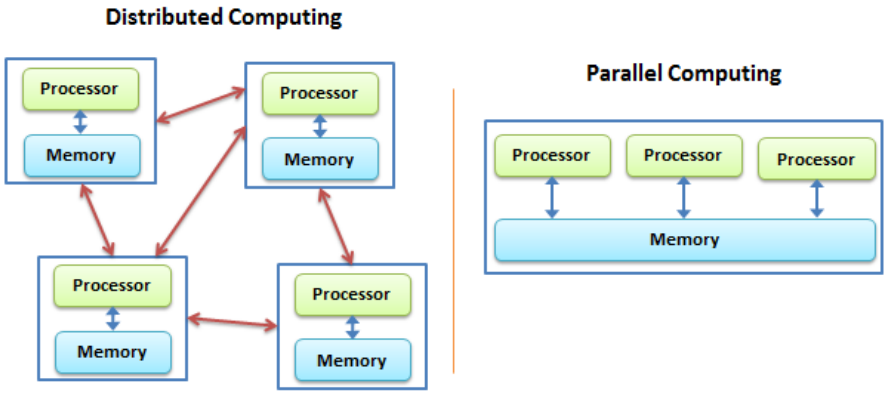

Aspect	Parallel Computing	Distributed Computing
Definition	Involves dividing a problem into sub-problems and solving them simultaneously on multiple processors in a single system.	Involves multiple autonomous systems (computers) working together to solve a problem, often connected over a network.
Architecture	Uses a single system with multiple processors (e.g., multi-core CPUs).	Uses multiple independent computers connected by a network (e.g., a cluster or the internet).
Memory	Typically uses shared memory.	Typically uses distributed memory (each system has its own memory).
Failure	Failure of one processor may affect the entire computation.	Failure of one system usually doesn't affect the entire system; other systems can continue to work.
Latency	Lower latency due to processors being within the same machine.	Higher latency due to network communication between distributed systems.
Scalability	Limited by the number of processors in a single system.	Highly scalable, as more machines can be added to the network.

Aspect	Parallel Computing	Distributed Computing
Applications	High-performance computing, simulations, data processing (e.g., supercomputers).	Cloud computing, distributed systems, blockchain (e.g., Google, Amazon).
Example	OpenMP, CUDA, MPI (Message Passing Interface) in tightly-coupled systems.	Hadoop, Spark, Kubernetes, Google Cloud.

▼ 🔰 **Q4. Explain Elements of Parallel Computing and Distributed Computing**

Elements of Parallel Computing -

1. **Task Breakdown:** Breaks tasks into smaller sub-tasks for simultaneous execution.
2. **Concurrency:** Multiple tasks run at the same time on different processors.
3. **Synchronization:** Ensures tasks are coordinated using locks, barriers, etc.
4. **Communication:** Tasks communicate via shared memory or message-passing.
5. **Load Balancing:** Evenly distributes tasks across processors.
6. **Speedup and Scalability:** Measures performance improvement and system efficiency as more processors are added.

Elements of Distributed Computing -

1. **Nodes:** Independent computers (or systems) that work together over a network.
2. **Concurrency:** Tasks run independently and often concurrently across nodes.
3. **Coordination:** Ensures tasks are synchronized using distributed locks or consensus protocols.
4. **Communication:** Nodes communicate through message-passing across a network.
5. **Fault Tolerance:** Uses redundancy and replication to handle node failures.
6. **Scalability:** Easily scales by adding more nodes to the network.

(Also Refer Unit-1-Q3)

▼ 🔰 **Q5. Write a Short Note on Full Virtualization. State Pros and Cons.**

- Full virtualization is a technique where a virtual machine (VM) simulates the complete hardware of a physical system, allowing unmodified operating systems and applications to run as if they were on real hardware.
- The hypervisor, or virtual machine monitor (VMM), creates and manages these VMs by emulating the underlying hardware resources like CPU, memory, and I/O devices.

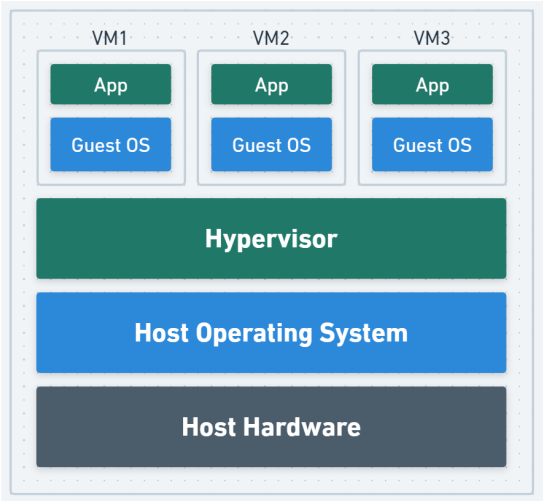

Pros -

1. **Isolation:** VMs are fully isolated from each other, improving security and stability.
2. **Compatibility:** Supports unmodified guest operating systems, making it flexible for different OS environments.
3. **Resource Utilization:** Maximizes hardware use by running multiple VMs on the same physical server.
4. **Portability:** VMs can be easily moved between servers without modification.

Cons -

1. **Overhead:** Emulating hardware introduces performance overhead, reducing efficiency compared to native execution.
2. **Complexity:** Managing full virtualized environments can be complex due to the need for advanced hypervisors.
3. **Resource Contention:** Multiple VMs on the same hardware may compete for resources, affecting performance.
4. **Hardware Requirements:** Requires powerful hardware to efficiently support multiple virtual machines.

▼ 🔰 **Q6. Explain The Machine Reference Model of Execution Virtualization.**

- When virtualizing an execution environment at different levels of the computation stack, a reference model is needed to define the interfaces between abstraction layers.
- These abstractions hide implementation details, allowing virtualization techniques to intercept calls and substitute layers.

- Clear separation between layers simplifies implementation by requiring only interface emulation and appropriate responses from the underlying layer.

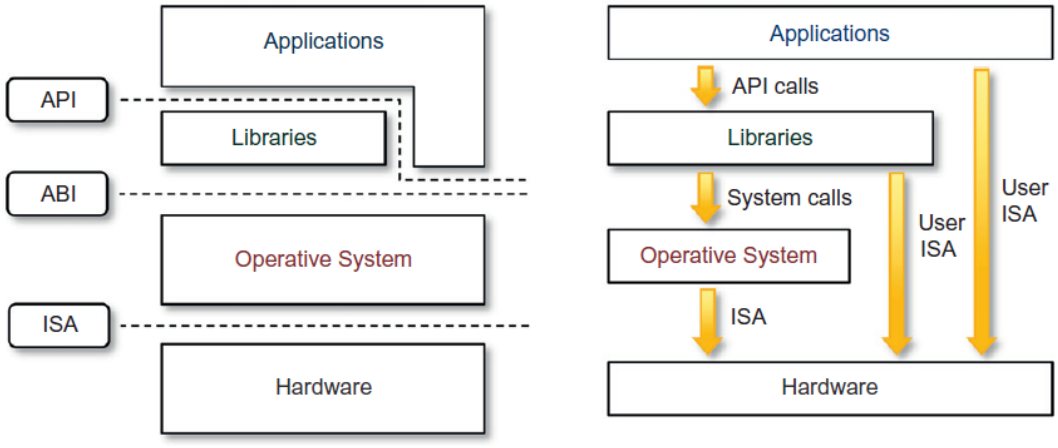

1. At the base layer, the model includes hardware architecture, specifically the **Instruction Set Architecture (ISA)**, which defines the processor's instruction set, registers, memory, and interrupt management, acting as the interface between software and hardware.
2. The **Application Binary Interface (ABI)** separates the operating system from applications and libraries, handling system calls and enabling portability across OSs using the same ABI. The **Application Programming Interface (API)**, the highest level of abstraction, interfaces applications with libraries and the core OS.
3. The CPU operates in two modes: **User Mode** (limited access to memory and no access to peripherals) and **Kernel Mode** (full access to memory and peripherals). Instructions are divided into **privileged** (for sensitive operations) and **non-privileged** (general operations without accessing shared resources).

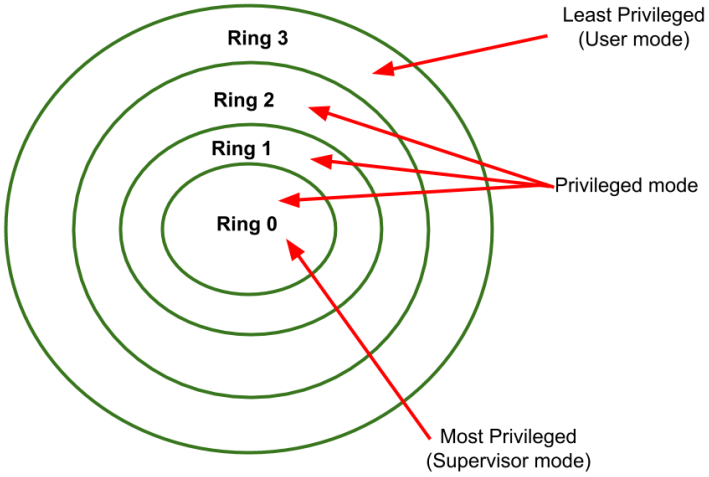

- In a hypervisor-managed environment, guest OS code typically runs in user mode to prevent direct OS access, but full isolation of guest OSs is not possible when using non-privileged instructions.

▼ 🔰 **Q7. Explain Hypervisor and Its Types.**

- A **Hypervisor**, also known as a **Virtual Machine Monitor (VMM)**, is software that creates and manages virtual machines (VMs) by enabling multiple operating systems to share a single physical hardware host.
- The hypervisor abstracts the hardware resources, such as CPU, memory, and storage, and allocates them to each VM, creating isolated environments for different operating systems to run concurrently on the same physical server.

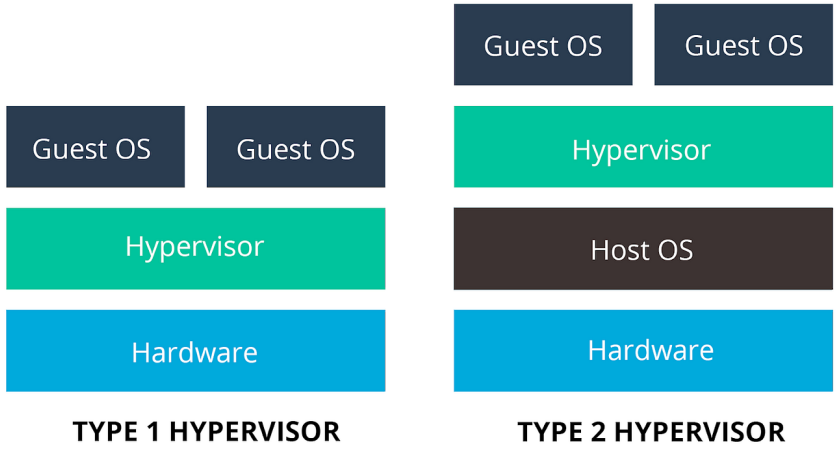

Types of Hypervisors -

1. **Type 1 Hypervisor (Bare-Metal Hypervisor) -**
   - Type 1 hypervisors run directly on the physical hardware of the host system, without the need for an underlying operating system. It acts as the operating system itself and manages the hardware resources and VMs.
   - **Examples:** Microsoft Azure, AWS, Xen.
   - **Pros -**
     1. Better Performance
     2. Lower Overhead (Load)
     3. Security and Stability
     4. Ideal for Enterprise
   - **Cons -**
     1. Complex Setup and Management
     2. Requires Specialized Hardware
     3. Less Accessible for Small-scale Users
2. **Type 2 Hypervisor (Hosted Hypervisor) -**

- Type 2 hypervisors run on top of an existing operating system (the host OS). It operates as an application within the OS, managing VMs.
- **Examples:** Oracle VirtualBox, VMware Workstation, Parallels Desktop.
- **Pros -**
  1. Easy to Install and Use
  2. Accessible for Small-scale Users
  3. Not Requires Specialized Hardware
- **Cons -**
  1. Lower Performance
  2. Higher Overhead (Load)
  3. Security Risks

▼ **Q8. Difference Between Full, Para, and Hardware-Assisted Virtualization.**

Aspect	Full Virtualization	Para-virtualization	Hardware-Assisted
Definition	Emulates complete hardware for unmodified guest OSs.	Requires modified guest OS that communicates with the hypervisor.	Utilizes CPU extensions to enhance virtualization performance.
Guest OS Modification	No modifications needed; can run unmodified OSs.	Requires modifications to the guest OS.	No modifications needed; supports unmodified OSs.
Performance	Higher overhead due to hardware emulation.	Lower overhead; better performance than full virtualization.	Better performance with minimal overhead.
Compatibility	High compatibility with various operating systems.	Limited compatibility; depends on guest OS modifications.	High compatibility, as it supports unmodified OSs.
Isolation	Strong isolation between VMs.	Good isolation, but less than full virtualization.	Strong isolation due to hardware-level support.
Use Cases	General-purpose virtualization, data centers.	High-performance environments, server farms.	Modern virtualization solutions, cloud services.
Example Technologies	VMware ESXi, Microsoft Hyper-V.	Xen.	Intel VT-x, AMD-V, VMware Workstation.

Unit 2 -

▼ 🔰 **Q1. Discuss Cloud Computing Roles and Boundaries.**

Cloud computing involves multiple stakeholders, each with specific roles and responsibilities that define the boundaries of interactions and services.

Roles -

1. **Cloud Service Provider (CSP) -**
   - **Description:** Organizations that offer cloud services, including Infrastructure as a Service (IaaS), Platform as a Service (PaaS), and Software as a Service (SaaS).

- **Responsibilities:**
    - Manage and maintain the cloud infrastructure.
    - Ensure availability, security, and compliance.
    - Provide technical support and customer service.

2. **Cloud Consumer -**
    - **Description:** Individuals or organizations that utilize cloud services provided by CSPs.
    - **Responsibilities:**
        - Select appropriate cloud services based on needs.
        - Manage resources and costs associated with cloud usage.
        - Ensure data security and compliance on their end.
3. **Cloud Auditor -**
    - **Description:** An independent entity that assesses the cloud environment for compliance, security, and performance.
    - **Responsibilities:**
        - Conduct audits to ensure CSPs meet security and compliance standards.
        - Review service agreements and SLAs.
        - Provide reports on cloud service performance and risks.
4. **Cloud Architect -**
    - **Description:** A professional responsible for designing and implementing cloud infrastructure and services.
    - **Responsibilities:**
        - Develop cloud strategies and roadmaps.
        - Design scalable and secure cloud architectures.
        - Ensure integration with existing systems and processes.
5. **Cloud Developer -**
    - **Description:** Software engineers who build applications and services to run in the cloud.
    - **Responsibilities:**
        - Develop and deploy cloud-native applications.
        - Utilize cloud services and APIs effectively.
        - Ensure applications are scalable and resilient.

Boundaries -

1. **Organizational Boundary -**

- An organizational boundary defines the physical and logical perimeter of an organization's IT resources.
- It encompasses the assets, infrastructure, and applications owned and managed by the organization.
- **Key Characteristics of Organizational Boundary -**
    1. **Physical:** Data centers, servers, network devices, and other physical infrastructure.
    2. **Logical:** Virtual networks, firewalls, and other security measures.
    3. **Ownership:** The organization owns and controls the resources within this boundary.
    4. **Governance:** The organization sets policies, procedures, and standards for managing these resources.

2. **Trust Boundary -**
    - A trust boundary is a logical perimeter that extends beyond the physical organizational boundary. It represents the extent to which an organization trusts external entities, such as cloud service providers, to handle its data and applications.
    - This boundary is often less defined and more complex, as it involves sharing control and responsibility with third-party providers.
    - **Key Characteristics of Trust Boundary -**
        1. **Shared Responsibility:** The organization and the cloud provider share responsibility for security, compliance, and data protection.
        2. **Shared Control:** The organization may have limited control over the underlying infrastructure, but it retains control over its data and applications.
        3. **Risk Assessment:** Organizations must carefully assess the risks associated with sharing data and applications with external providers.
        4. **Contractual Agreements:** Clear contractual agreements are essential to define the terms of the relationship and the responsibilities of each party.

▼ 🔰 **Q2. Explain Cloud Computing Delivery and Deployment Models.**

Cloud computing is characterized by various delivery and deployment models that define how services are provided and utilized. These models cater to different needs and use cases, offering flexibility and scalability.

Cloud Delivery Models -

1. **Infrastructure as a Service (IaaS) -**
    - **Description:** Provides virtualized computing resources over the internet. Users can rent virtual machines, storage, and networks without managing physical hardware.
    - **Use Cases:** Hosting websites, running applications, and data storage.

- **Examples:** Amazon EC2, Microsoft Azure Virtual Machines, Google Cloud Compute Engine.

2. **Platform as a Service (PaaS) -**
   - **Description:** Offers a platform allowing developers to build, deploy, and manage applications without dealing with the underlying infrastructure.
   - **Use Cases:** Application development, testing, and deployment.
   - **Examples:** Google App Engine, Microsoft Azure App Service, Heroku.
3. **Software as a Service (SaaS) -**
   - **Description:** Delivers software applications over the internet on a subscription basis. Users access applications via a web browser without needing installation.
   - **Use Cases:** Email, collaboration tools, and customer relationship management (CRM).
   - **Examples:** Google Workspace, Microsoft 365, Salesforce.

Cloud Deployment Models -

1. **Public Cloud -**
   - **Description:** Services are offered over the public internet and shared among multiple organizations. Resources are owned and managed by third-party providers.
   - **Pros:** Cost-effective, scalable, and easily accessible.
   - **Cons:** Less control over security and compliance.
   - **Examples:** AWS, Microsoft Azure, Google Cloud Platform.
2. **Private Cloud -**
   - **Description:** Dedicated cloud infrastructure for a single organization. It can be hosted on-premises or by a third-party provider, offering greater control and customization.
   - **Pros:** Enhanced security, compliance, and control over resources.
   - **Cons:** Higher costs and requires management.
   - **Examples:** VMware Private Cloud, OpenStack.
3. **Hybrid Cloud -**
   - **Description:** Combines public and private clouds, allowing data and applications to be shared between them. Organizations can leverage the benefits of both models.
   - **Pros:** Flexibility, scalability, and improved security for sensitive data.
   - **Cons:** More complex management and potential integration challenges.
   - **Examples:** AWS Outposts, Microsoft Azure Stack.

▼ 🔰 **Q3. Explain Open Challenges of Cloud Computing and Economics Of the Cloud.**

Open Challenges of Cloud Computing -

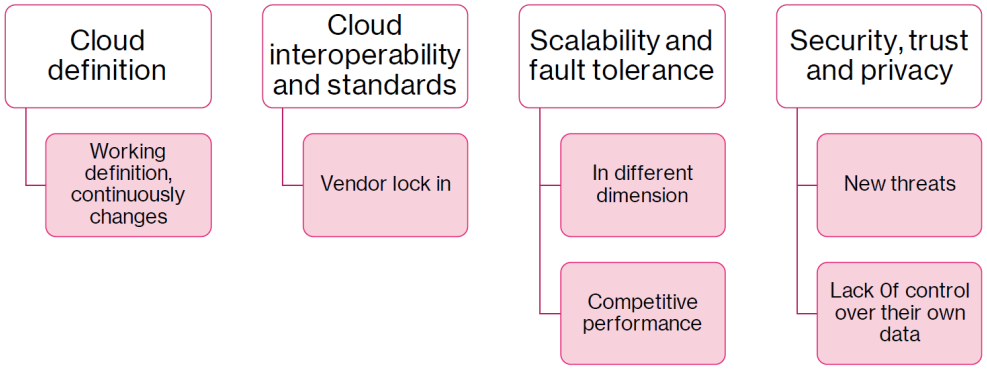

1. **Security and Privacy -**
   - **Challenge:** Ensuring data security and privacy remains a major concern, especially with sensitive information stored in the cloud.
   - **Considerations:** Risks include data breaches, unauthorized access, and compliance with regulations like GDPR.
2. **Data Governance and Compliance -**
   - **Challenge:** Organizations must navigate complex regulatory requirements regarding data management and sovereignty.
   - **Considerations:** Ensuring compliance with laws across different jurisdictions can be difficult, especially in multi-cloud environments.
3. **Vendor Lock-In -**
   - **Challenge:** Difficulty in migrating applications and data between different cloud providers can lead to dependence on a single vendor.
   - **Considerations:** Organizations may face challenges in terms of costs and compatibility when switching providers.
4. **Interoperability -**
   - **Challenge:** The lack of standardization among cloud services can hinder integration between different platforms.
   - **Considerations:** Ensuring seamless communication and data exchange between different cloud environments can be complex.
5. **Performance and Reliability -**
   - **Challenge:** Ensuring consistent performance and reliability of cloud services is crucial for user satisfaction.

- **Considerations:** Issues such as latency, downtime, and resource allocation can affect application performance.

6. **Scalability and Elasticity -**
   - **Challenge:** While cloud platforms offer scalability, planning for sudden demand spikes can be challenging.
   - **Considerations:** Organizations must ensure their architecture can handle variable workloads efficiently.

Economics of the Cloud -

1. **Cost Savings -**
   - Reduces capital expenditures by eliminating the need for physical infrastructure.
2. **Pay-as-You-Go Pricing -**
   - Flexible consumption-based model allows organizations to scale resources as needed.
3. **Resource Utilization -**
   - Efficient pooling of resources among multiple clients reduces waste.
4. **Operational Efficiency -**
   - Easy access to advanced technologies enhances overall efficiency.
5. **Financial Models -**
   - Offers flexibility in budgeting through CapEx and OpEx options.
6. **Risk Management -**
   - Mitigates financial risks with flexible contracts and pay-per-use models.

▼ 🔰 Q4. Explain Cloud Security: Threat Agents, Security Threats, and Additional Considerations.

Cloud security is critical due to the unique vulnerabilities and risks associated with cloud computing environments. Understanding the various threat agents and security threats, along with additional considerations, is essential for effective cloud security management.

Threat Agents -

1. **Anonymous Attacker -**
   - An external attacker who hides their identity, using various techniques to exploit vulnerabilities without revealing themselves.
2. **Malicious Service Agent -**
   - A compromised or rogue service within the cloud that intentionally harms users or steals data, often posing as a legitimate service.
3. **Trusted Attacker -**

- An attacker who has legitimate access to the system (e.g., an employee) and uses that trust to exploit resources or data for malicious purposes.

4. **Malicious Insider -**
    - An individual within the organization (like an employee or contractor) who misuses their access to sensitive information or resources, often for personal gain.

Security Threats -

1. **Traffic Eavesdropping -**
    - Intercepting data transmitted over networks, allowing attackers to access sensitive information without detection.
2. **Malicious Intermediary -**
    - A compromised entity that intercepts and potentially alters communication between users and cloud services, often to steal data or credentials.
3. **Denial of Service (DoS) -**
    - Attacks that overwhelm cloud services with excessive requests, making them unavailable to legitimate users.
4. **Insufficient Authorization -**
    - Weak access controls that allow unauthorized users to access resources or data they shouldn't, leading to potential data breaches.
5. **Virtualization Attack -**
    - Exploiting vulnerabilities in the hypervisor or virtual machines, allowing attackers to gain control over multiple VMs or the host system.
6. **Overlapping Trust Boundaries -**
    - Situations where different security zones overlap, creating vulnerabilities due to unclear trust relationships between systems and users.

Additional Considerations -

1. **Flawed Implementations -**
    - Errors in configuration or design of cloud services that can create security vulnerabilities, leading to potential data breaches or system failures.
2. **Security Policy Disparity -**
    - Differences in security policies between cloud providers and customers can lead to gaps in protection and compliance, risking data security.
3. **Contracts -**
    - Service Level Agreements (SLAs) and contracts outline the responsibilities, expectations, and security measures between cloud providers and consumers, impacting accountability

and risk.

4. **Risk Management -**
    - The process of identifying, assessing, and mitigating risks associated with cloud computing to protect data and ensure compliance, including regular audits and updates to security measures.

▼ 🔰 **Q5. Write a Short Note on Amazon Web Services, Google App Engine, and Microsoft Azure.**

Amazon Web Services (AWS) -

- **Overview:** AWS is a comprehensive cloud computing platform offered by Amazon, providing a wide range of services including computing power, storage options, and machine learning.
- **Key Features:**
    - Scalability: Offers scalable resources to handle varying workloads.
    - Extensive Services: Includes over 200 fully featured services, including EC2 (Elastic Compute Cloud), S3 (Simple Storage Service), and Lambda for serverless computing.
    - Global Reach: Data centers across multiple regions worldwide enable low-latency access and redundancy.
- **Use Cases:** Suitable for startups to enterprises, supporting web hosting, data analysis, and application development.

Google App Engine -

- **Overview:** Google App Engine is a platform-as-a-service (PaaS) offered by Google that allows developers to build and host applications on Google's infrastructure.
- **Key Features:**
    - Automatic Scaling: Automatically scales applications based on traffic without requiring manual intervention.
    - Built-in Services: Provides built-in services like NoSQL databases, user authentication, and monitoring tools.
    - Integration: Seamlessly integrates with other Google Cloud services, such as BigQuery and Cloud Storage.
- **Use Cases:** Ideal for web applications and APIs, especially for developers looking to leverage Google's AI and machine learning capabilities.

Microsoft Azure -

- **Overview:** Azure is Microsoft's cloud computing platform that offers a range of services, including analytics, storage, and networking, with a strong emphasis on enterprise solutions.
- **Key Features:**

    - Hybrid Capabilities: Supports hybrid cloud solutions, allowing organizations to integrate on-premises data centers with the cloud.
    - Extensive Tooling: Offers development tools like Visual Studio and DevOps services for streamlined application development and deployment.
    - Comprehensive Security: Built-in security features and compliance certifications to meet enterprise needs.
- **Use Cases:** Suitable for businesses of all sizes, especially for enterprises looking for a flexible cloud solution that integrates with Microsoft products.